

API For Your LLM

A **FastAPI REST API** that wraps a local **Ollama LLM** (Mistral model) with **API key authentication** and a **credit-based usage system**.

· Features

- **·API Key Auth** · Header-based authentication via `x-api-key`
 - **·Credit System** · Usage tracking with credit deduction per request
 - **·Local LLM** · Powered by Ollama (Mistral model)
 - **·FastAPI** · High-performance async REST endpoints
-

.. Architecture

```
Client · POST /generate (x-api-key) · FastAPI · verify_api_key() · deduct
credit · ollama.chat() · Response
```

· Getting Started

Prerequisites

- Python 3.10+, Ollama running, `ollama pull mistral`

```
cd API-For-Your-LLM
pip install -r requirements.txt
uvicorn main:app --reload
```

Test

```
curl -X POST http://localhost:8000/generate \
  -H "x-api-key: your-key" \
  -H "Content-Type: application/json" \
  -d '{"prompt": "Explain quantum computing"}'
```

· API Reference

POST /generate

Header	Required	Description	----- ----- -----	x-api-key	Yes	API key for authentication
--------	----------	-------------	-------------------	-----------	-----	----------------------------

Request: `{"prompt": "Your question"}` **Response:** `{"response": "Generated text..."}`

· Logic Flow

- 1 Client sends POST with prompt and API key

- 2 `verify_api_key()` validates the key
 - 3 Credit check and deduction
 - 4 Forward prompt to Ollama's Mistral model
 - 5 Return generated text as JSON
-

· Dependencies

`fastapi, uvicorn, ollama, python-dotenv, requests`

· License

Educational project · use freely for learning and reference.